



**ADDIS ABABA SCIENCE AND TECHNOLOGY
UNIVERSITY
COLLEGE OF ENGINEERING
DEPARTMENT OF SOFTWARE ENGINEERING**

**SOFTWARE COMPONENT DESIGN
Python ML Churn Prediction Model**

Section C

Name	ID
Elbetel Shineda	ETS0406/13
Leul Mintesinot	ETS0766/13
Mihret Agegnehu	ETS0874/13
Mikiyas Bedasa	ETS0888/13
Natanim Ashenafi	ETS0979/13

Submitted to: Mr. Gizate Desalgn
Submission Date: December 18, 2024

Machine Learning Predicting Customer Churn	3
Introduction	3
Project Setup	3
Tools and Libraries Used	3
Dataset	3
Data Preprocessing	3
1. Handling Missing Values	3
2. Encoding Categorical Variables	4
3. Scaling Numerical Features	4
Exploratory Data Analysis (EDA)	4
1. Distribution of Churn	4
2. Correlation Analysis	4
3. Feature Relationships	5
Model Development	5
1. Splitting the Data	5
2. Model Selection	5
3. Hyperparameter Tuning	5
Model Evaluation	6
Metrics Used	6
Key Findings	6
Code Execution Instructions	6
Prerequisites	6
Steps	7
Conclusion	7

Machine Learning Predicting Customer Churn

Introduction

This document provides a detailed explanation of the implementation, methodology, and logic behind the Python Machine Learning code for predicting customer churn. The code uses the Telco Customer Churn dataset to build a machine learning model that identifies customers likely to discontinue services.

Project Setup

Tools and Libraries Used

1. **pandas**: For data manipulation and preprocessing.
2. **numpy**: For numerical computations.
3. **matplotlib** and **seaborn**: For data visualization.
4. **scikit-learn**: For machine learning models and evaluation metrics.

Dataset

- **Source**: Telco Customer Churn dataset from Kaggle.
- **Size**: 7043 rows and 21 columns.
- **Target Variable**: **Churn** (Binary: Yes/No).
- **Features**: Includes customer demographics, service usage patterns, and account information.

Data Preprocessing

1. Handling Missing Values

- Columns like **TotalCharges** contained missing values due to empty strings.

Solution: Convert empty strings to NaN and replace NaN values with the mean of the column.

```
df['TotalCharges'] = pd.to_numeric(df['TotalCharges'], errors='coerce')
```

- `df['TotalCharges'].fillna(df['TotalCharges'].mean(), inplace=True)`

2. Encoding Categorical Variables

Categorical variables (e.g., **gender**, **Partner**, **Dependents**) were converted to numeric values using **Label Encoding** or **One-Hot Encoding** depending on their nature.

```
from sklearn.preprocessing import LabelEncoder  
le = LabelEncoder()
```

- `df['gender'] = le.fit_transform(df['gender'])`

3. Scaling Numerical Features

Features like **tenure** and **MonthlyCharges** were normalized using **Min-Max Scaling** to improve model performance.

```
from sklearn.preprocessing import MinMaxScaler  
scaler = MinMaxScaler()
```

- `df[['tenure', 'MonthlyCharges', 'TotalCharges']] = scaler.fit_transform(df[['tenure', 'MonthlyCharges', 'TotalCharges']])`

Exploratory Data Analysis (EDA)

1. Distribution of Churn

- The dataset has an approximate churn rate of **26.5%**.
- Visualized using a bar plot:
`sns.countplot(df['Churn'])`

2. Correlation Analysis

- Heatmaps were generated to identify correlations between features and the target variable **Churn**.

```
sns.heatmap(df.corr(), annot=True, cmap='coolwarm')
```

3. Feature Relationships

- Histograms and box plots were created to analyze numerical features like **tenure** and **MonthlyCharges**.

Model Development

1. Splitting the Data

The dataset was split into **training** (80%) and **test** (20%) sets.

```
from sklearn.model_selection import train_test_split
```

```
X = df.drop('Churn', axis=1)
```

```
y = df['Churn']
```

- `X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)`

2. Model Selection

- **Algorithms Tested:**
 - Logistic Regression
 - Random Forest
 - Gradient Boosting
 - K-Nearest Neighbors (KNN)
 - Support Vector Machines (SVM)
- KNN performed the best with an accuracy of **98.2%**.

3. Hyperparameter Tuning

Grid search was used to optimize hyperparameters for each model.

```
from sklearn.model_selection import GridSearchCV
```

```
param_grid = {'n_neighbors': [3, 5, 7], 'weights': ['uniform', 'distance']}  
grid = GridSearchCV(KNeighborsClassifier(), param_grid, cv=5)
```

- `grid.fit(X_train, y_train)`

Model Evaluation

Metrics Used

1. **Accuracy:** Proportion of correctly classified samples.
2. **Precision:** Proportion of positive predictions that are correct.
3. **Recall:** Proportion of actual positives that are correctly identified.
4. **F1-Score:** Harmonic mean of precision and recall.

Example of evaluation:

```
from sklearn.metrics import classification_report  
y_pred = grid.best_estimator_.predict(X_test)  
print(classification_report(y_test, y_pred))
```

Key Findings

- **Important Features:** `tenure`, `MonthlyCharges`, and `Contract` had the most influence on churn predictions.
- **Best Model:** KNN achieved the highest accuracy of **98.2%**.

Code Execution Instructions

Prerequisites

- Python 3.8 or higher.
- Libraries: pandas, numpy, matplotlib, seaborn, scikit-learn.

Steps

1. Clone the repository:

```
git clone https://github.com/mihretgold/churn_prediction_group2.git
```

2. Navigate to the project directory:

```
cd churn-prediction
```

3. Run the notebook:

```
jupyter notebook churn_prediction.ipynb
```

Conclusion

This documentation provides an in-depth explanation of the code implementation, from data preprocessing to model evaluation. The KNN model emerged as the best-performing algorithm, enabling accurate customer churn predictions to help businesses retain customers and minimize revenue loss.